Install Greenplum on Docker and Connect via SqlWorkbench

Client	N/A
Task Definition	How to install and run GreenPlum database on Docker
Prepared By	Ahsanul Hadi
Date	24.08.2017
Checked By	

Note:

- This was done on Mac OS Sierra v: 10.12.5
- Docker v: 17.06.1-ce-mac24 (18950)
- Follow the below steps sequentially

Install Docker

https://docs.docker.com/engine/installation/

If you are not familiar with Docker, then here's a great tutorial that will help you to understand - what docker is, how it works and use cases.

http://takacsmark.com/getting-started-with-docker-in-your-project-step-by-step-tutorial/

Install GreenPlum

Source of Docker image for this installation: https://hub.docker.com/r/pivotaldata/gpdb-base/

Open up your Terminal and execute below commands -

Code block

```
$ cd ~
$ docker pull pivotaldata/gpdb-base
```

Check list of installed images

```
$ docker images
```

Output:

```
cbigadmins-MacBook-Pro:~ ahsanulhadi$ docker images
REPOSITORY                TAG                 IMAGE ID            CREATED             SIZE
hello-world               latest              1815c82652c0        2 months ago        1.84kB
pivotaldata/gpdb-base     latest              bfe4e63b8e81        18 months ago       1.17GB
```

Run the Docker image

Syntax: docker run -i -d -p 5432:5432 <image>

Arguments: **-p** = port (First port is for Host and the second one is for the container where it will be pointed to / **-d** = run as daemon / **-i** = interactive mode / **--name** = add name

For detail explanation on these parameters, check here: https://docs.docker.com/engine/reference/commandline/run/

```
$ docker run --name gpdb_dev -i -d -p 5432:5432 bfe4e63b8e81
```

output will be a similar id: ed1830a37eded65687d6d3fed2f745041989040a4749890e45aa6d250c81050d

But if we **want to ssh into the container then run this below command**. (Add -p 2022:22)

Important: if you have already ran it then Stop the container first and remove it. (see last section)

```
$ docker run --name gpdb_dev -i -d -p 5432:5432 -p 2022:22 bfe4e63b8e81
```

Output:

```
cbigadmins-MacBook-Pro:~ ahsanulhadi$
cbigadmins-MacBook-Pro:~ ahsanulhadi$ docker run --name gpdb_dev -i -d -p 5432:5432 -p 2022:22 bfe4e63b8e81
93e0ed0088ca5d5ac4aded1257325622dcc78f188437bf33cb6e051fae755473
```

You can check with 'docker ps' command to see whether the container is running or not.

```
$ docker ps -a
```

Output:

```
cbigadmins-MacBook-Pro:~ ahsanulhadi$ docker ps -a
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS                      PORTS                                        NAMES
93e0ed0088ca        bfe4e63b8e81        "/bin/sh -c 'echo ..."   About a minute ago  Up About a minute           0.0.0.0:5432->5432/tcp, 0.0.0.0:2022->22/tcp   gpdb_dev
7959879d8cfe        hello-world         "/hello"                 2 days ago          Exited (0) 2 days ago                                                    cocky_stonebraker
```

SSH into GreenPlum Container

By now we have the GreenPlum database container running on 0.0.0.0:5432. Now log in via SSH and connect to psql (local version) or pgadmin works as well.

User: **gpadmin** Password: **pivotal** OR **root**/pivotal

```
$ ssh gpadmin@0.0.0.0 -p 2022
```

Output:

```
cbigadmins-MacBook-Pro:~ ahsanulhadi$ ssh gpadmin@0.0.0.0 -p 2022
The authenticity of host '[0.0.0.0]:2022 ([0.0.0.0]:2022)' can't be established.
RSA key fingerprint is SHA256:RBXjRzXPVrY3tWhQD+QtyN6hTgF2anVau8yDsURWx9c.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '[0.0.0.0]:2022' (RSA) to the list of known hosts.
gpadmin@0.0.0.0's password:
-bash-4.1$
-bash-4.1$
```

Copy files to GreenPlum Server/Container

If you need to copy files into Greenplum server from your local machine, here is the command:

```
$ scp -P 2022 my_local_data_file.csv
gpadmin@0.0.0.0:/home/gpadmin/destination_dir/
```

Connect via psql

```
cbigadmins-MacBook-Pro:~ ahsanulhadi$ ssh gpadmin@0.0.0.0 -p 2022
gpadmin@0.0.0.0's password:
Last login: Thu Aug 24 00:30:24 2017 from 172.17.0.1
-bash-4.1$
-bash-4.1$ psql
psql (8.2.15)
Type "help" for help.

gpadmin=# \l+
                                    List of databases
   Name    |  Owner   | Encoding |  Access privileges  | Size  | Tablespace |        Description
-----------+----------+----------+---------------------+-------+------------+---------------------------
 gpadmin   | gpadmin | UTF8     |                     | 29 MB | pg_default |
 postgres  | gpadmin | UTF8     |                     | 29 MB | pg_default |
 template0 | gpadmin | UTF8     | =c/gpadmin          | 27 MB | pg_default |
                                : gpadmin=CTc/gpadmin
 template1 | gpadmin | UTF8     | =c/gpadmin          | 29 MB | pg_default | Default template database
                                : gpadmin=CTc/gpadmin
(4 rows)

gpadmin=# \q
-bash-4.1$ exit
logout
Connection to 0.0.0.0 closed.
cbigadmins-MacBook-Pro:~ ahsanulhadi$
```

Stop the container

```
$ docker stop gpdb_dev
$ docker start gpdb_dev # if required
```

Output:

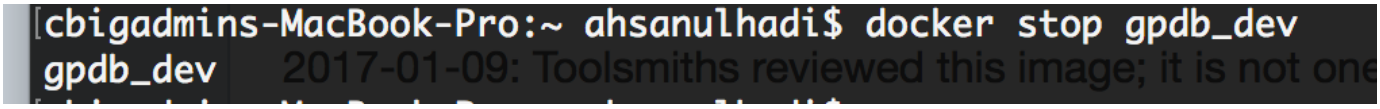

Remove the container

```
$ docker rm gpdb_dev
```

Output:

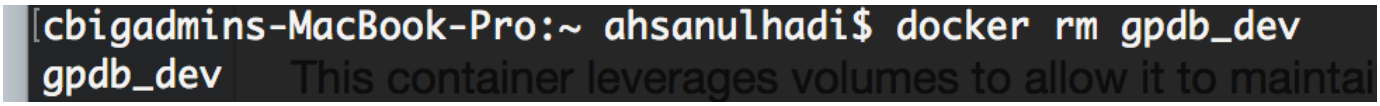

Download and configure SQL Workbench

Prerequisite:

- Go to http://www.sql-workbench.net/ . and download .dmg file.
- Go to https://jdbc.postgresql.org/download.html . and download appropriate JDBC driver. (check your Java version. Command: $ **java -version**)
- If you don't have Java installed then go to https://www.java.com/en/download/help/index_installing.xml

Install SQL Workbench and configure the driver. Open the application and go to 'manage driver'. Or you can go from 'File Manage Drivers'. See image below:

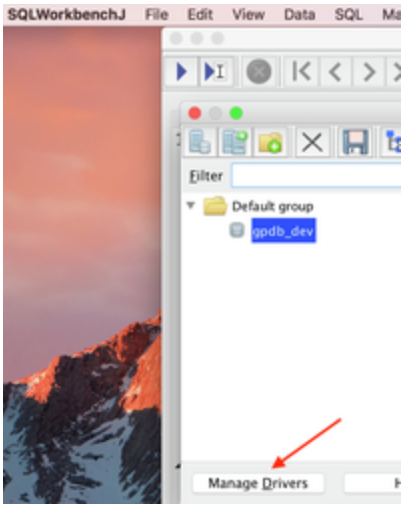

Update the driver file location for Postgres Driver. Select the driver that we have downloaded in previous step.

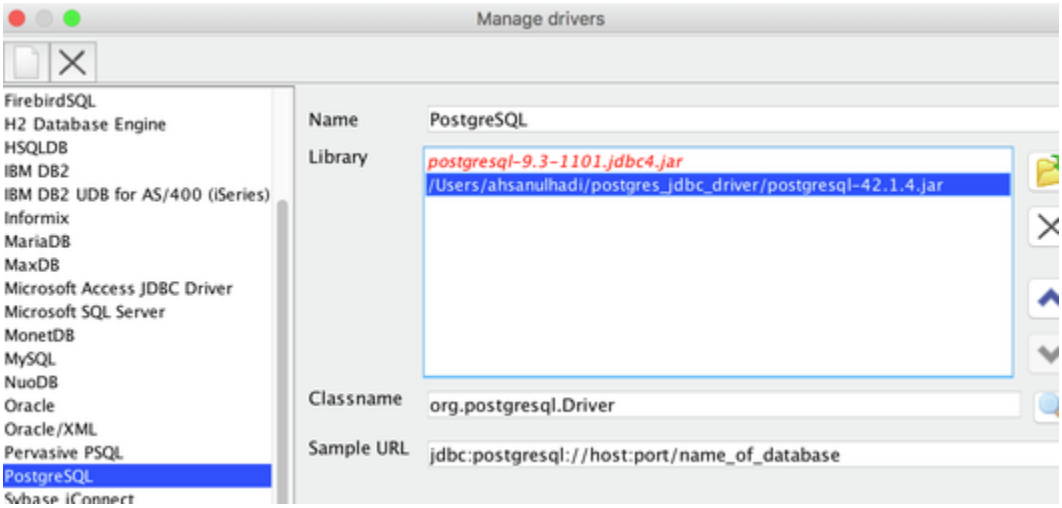

Now, create new Database co nnection profile.

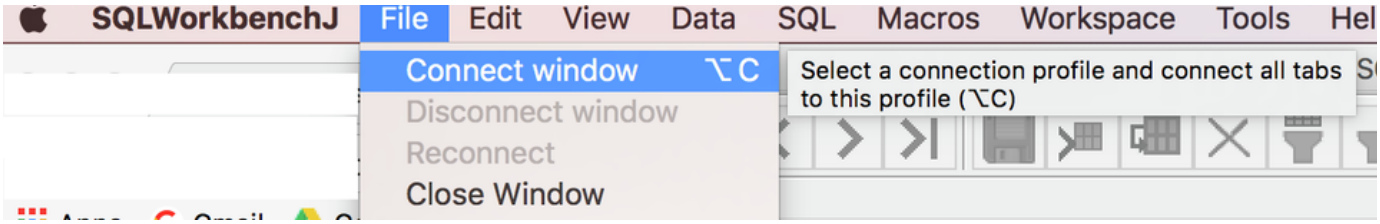

Add the database connection info and save the profile.

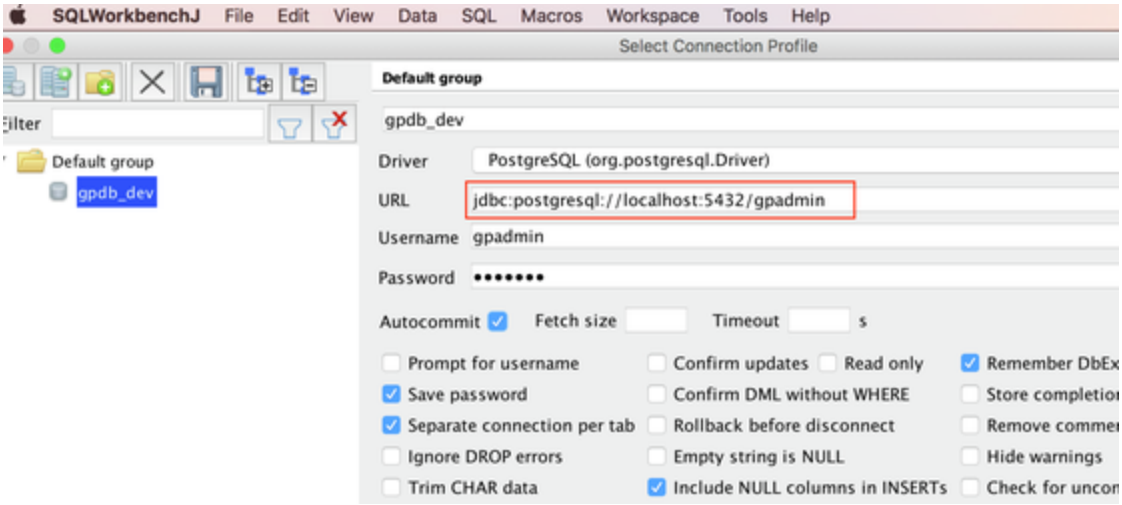
SQLWorkbenchJ File Edit View Data SQL Macros Workspace Tools Help
Select Connection Profile
Default group
gpdb_dev
Driver PostgreSQL (org.postgresql.Driver)
URL jdbc:postgresql://localhost:5432/gpadmin
Username gpadmin
Password •••••••
Autocommit Fetch size Timeout s
Prompt for username
Save password
Separate connection per tab
Ignore DROP errors
Trim CHAR data
Confirm updates Read only
Confirm DML without WHERE
Rollback before disconnect
Empty string is NULL
Include NULL columns in INSERTs
Remember DbEx
Store completio
Remove comme
Hide warnings
Check for uncon

And connect 🙂

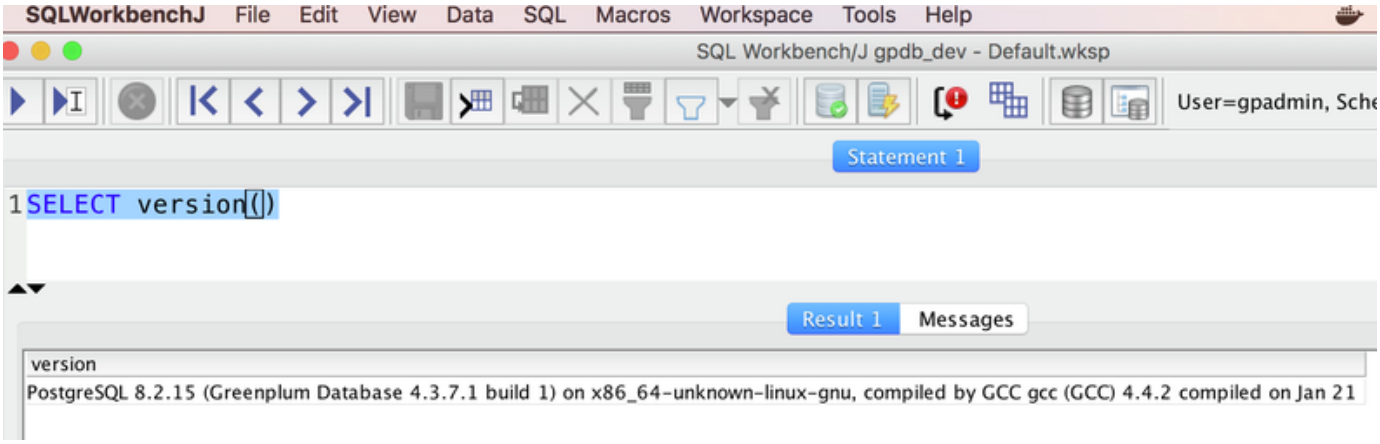
SQLWorkbenchJ File Edit View Data SQL Macros Workspace Tools Help
SQL Workbench/J gpdb_dev - Default.wksp
User=gpadmin, Sche
Statement 1
1 SELECT version()
Result 1 Messages
version
PostgreSQL 8.2.15 (Greenplum Database 4.3.7.1 build 1) on x86_64-unknown-linux-gnu, compiled by GCC gcc (GCC) 4.4.2 compiled on Jan 21